

# Project report — Old-car accident classification

## Executive summary

I converted a used-cars dataset into a machine-learning-ready table, engineered features (including missingness flags and groupings for rare categories), encoded categorical variables (one-hot and frequency), split data into train/val/test (stratified), trained an XGBoost classifier using SageMaker built-in XGBoost, and evaluated it locally using the trained model artifact. The model achieves  $AUC \approx 0.7014$ . With a threshold chosen for high recall ( $\approx 0.3471$ ), it reaches recall  $\approx 0.80$ , precision  $\approx 0.35$ ,  $F1 \approx 0.488$ . I iterated on preprocessing and thresholding rather than deploying an endpoint.

## 1. Environment & tooling

- Local Jupyter for EDA, preprocessing and feature engineering.
- SageMaker (console + notebook instance) for training with built-in XGBoost algorithm; S3 for data + model artifacts.

## 2. Data cleaning & parsing

I inspected columns and converted text fields into proper numeric/categorical formats so models can use them.

### Price, mileage, model\_year, engine

- **Price:** removed currency symbols and commas, then converted to numeric. This makes price comparable across rows.
- **Mileage:** stripped non-digits (commas, text like "km") then converted to numeric.
- **Model year:** forced to numeric.
- **Engine size:** parsed the engine text (e.g., 2.0L) and extracted the numeric part as engine\_l (e.g., 2.0).
- I used **errors='coerce'** on conversions so bad values become NaN and are handled explicitly rather than causing crashes.

### 3. Target creation and missing-target handling

- I inspected accident values and created accident\_label with 0 for “None reported” and 1 for “At least 1 accident or damage reported.”
- I dropped rows with missing accident labels because a supervised classifier needs labeled examples.

### 4. Missingness as a feature

- I created binary flags that record whether a column value was missing: e.g., clean\_title\_missing, fuel\_type\_missing, engine\_l\_missing (1 = missing, 0 = present).
- I checked whether missingness correlated with the accident label (cross-tabulations). Missingness often had different accident rates than non-missing rows.

### 5. Imputation (how I filled missing values)

- Categorical columns (clean\_title, fuel\_type, etc.): I filled missing entries with the string "Unknown". I kept side-by-side \_missing flags so the model knows which "Unknown"s were originally missing.
- Numeric columns (engine\_l): I filled missing values with the median of the column.

## 6. Categorical handling — rare groups & encoding

I divided categorical encoding into two strategies to keep the feature set manageable and meaningful.

### 6.1 Rare-grouping

- For very high-cardinality columns (especially model and ext\_col), I combined categories that appear fewer than 10 times into a single bucket "Other".
- On this dataset:
  - model had many unique values → 1830 rare models were grouped into model\_grp = "Other".
  - ext\_col (exterior color) had many values → 298 rare colors were grouped into ext\_col\_grp = "Other".
  - Columns with fewer uniques (e.g., brand with 57 uniques) were left as-is.

### 6.2 Encoding after grouping

- **One-hot encoding (create binary columns for each category) for low-cardinality columns:** I applied this to fuel\_type and clean\_title. These have a small number of distinct values and one-hot is straightforward.
- **Frequency encoding for higher-cardinality columns:** I converted brand, model\_grp, transmission, ext\_col\_grp, int\_col into numeric features by replacing each category with its frequency (the fraction of rows having that value). Frequency encoding keeps representation compact and often conveys useful information about common vs rare categories.

## 7. Train/val/test split and class imbalance

- I split the dataset into train (70%), validation (15%) and test (15%) using stratified split to keep accident proportions consistent between train/test.
- I computed scale\_pos\_weight = (#negatives / #positives) from the training set ( $\approx 2.95$ ) to inform XGBoost that accidents are the minority class.

## 8. SageMaker Training

- I uploaded the preprocessed training and validation CSV files to S3 under my bucket/prefix.
- I launched a small SageMaker Notebook instance for the console and ran one training job using the built-in XGBoost container. I used a modest instance type (e.g., ml.t3.medium) and made sure to stop the notebook after launching training.
- I used XGBoost with objective=binary:logistic and eval\_metric=auc. I ran a baseline job, examined logs, then submitted a tuned job with adjusted hyperparameters (lower learning rate, smaller max depth, modest regularization, and scale\_pos\_weight from the training set).
- The training job completed; the model artifact was saved to S3.

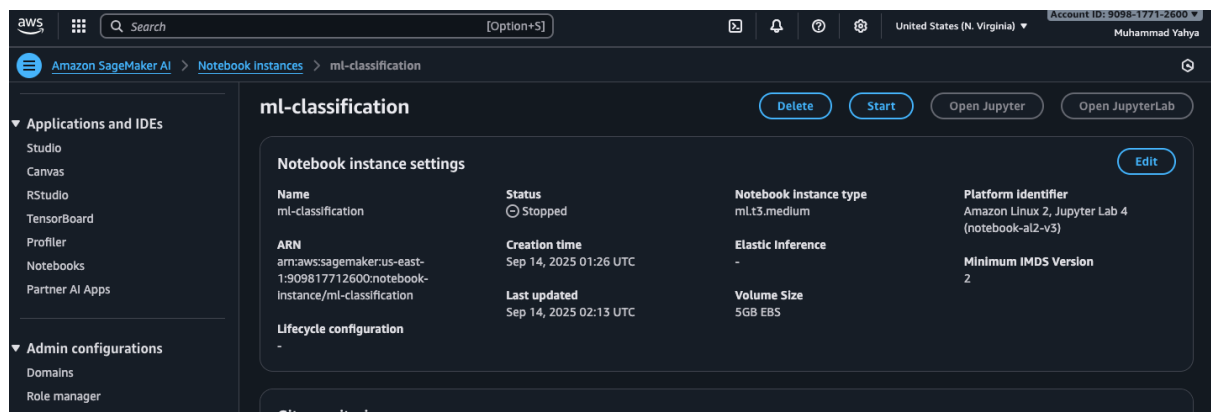


Figure 1 SageMaker Notebook Instance

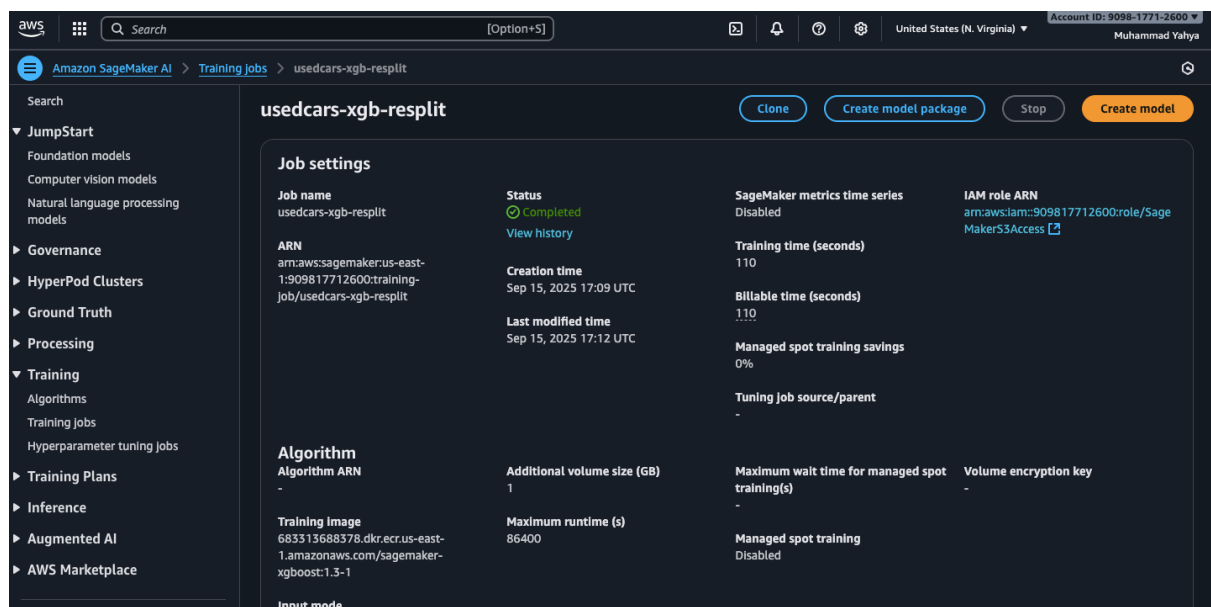


Figure 2 SageMaker Training Job

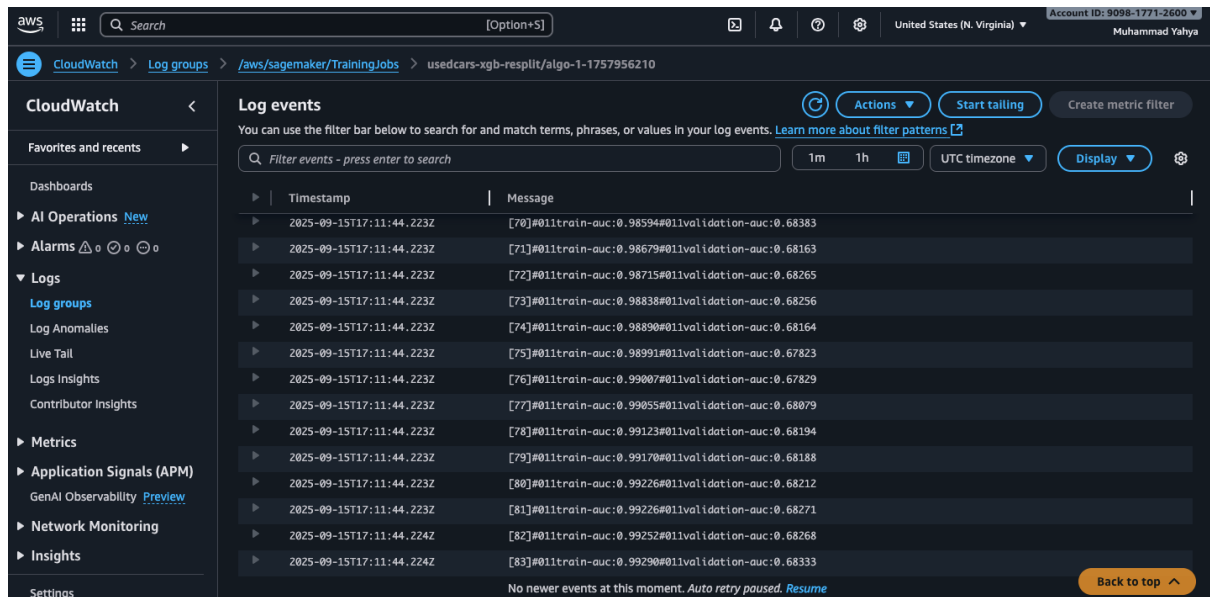


Figure 3 CloudWatch Training Job Logs

## 9. Model artifact and local evaluation

- I downloaded model.tar.gz from S3, extracted it locally, and loaded the XGBoost model for evaluation.
- I loaded the preprocessed test CSV (that contains the true accident\_label) and prepared the feature matrix matching the final\_features ordering.
- I calculated predicted probabilities, then evaluated performance.
- Final evaluation results (key :
  - Test AUC: 0.7014 — the model ranks positives better than random.
  - With threshold chosen to satisfy  $\text{recall} \geq 0.80$  (threshold  $\approx 0.3471$ ):
    - Precision: 0.35
    - Recall: 0.8041
    - F1: 0.4877

- Confusion matrix:
  - True negatives: 216
  - False positives: 221
  - False negatives: 29
  - True positives: 119

## 10. Conclusion

In conclusion, I successfully built and evaluated a machine learning model to predict the likelihood of a used car having an accident history. The data required significant cleaning, feature engineering (including handling missing values, creating missingness flags, and encoding categorical variables), and careful preprocessing before training. I then trained and tested the model using SageMaker and later evaluated it locally. The final model achieved an AUC of 0.70, with good recall (0.80) for accident cases, meaning it was effective at identifying cars with accident history, though it came at the cost of lower precision (0.35). Overall, the project demonstrates a strong foundation for accident risk prediction and provides a scalable framework that can be further improved with more balanced data and feature refinements.